

Lab 9. Running Elastic Stack in Production



We will start with Elasticsearch and then move on to other components. There are various ways to run Elasticsearch in production. There may be various factors that influence your decision on how you should deploy. We will cover the following topics to help you take your next Elastic Stack project to production:

- Hosting Elastic Stack on a managed cloud
- Hosting Elastic Stack on your own, that is, self-hosting
- Backing up and restoring
- Setting up index aliases
- Setting up index templates
- Modeling time series data

Let's first understand how we can go about taking Elastic Stack to production with one of the managed cloud providers. This option requires a minimum amount of work to set up a production-ready cluster.

Hosting Elastic Stack on a managed cloud

Cloud providers make the process of setting up a production-ready cluster much easier. As a user, we don't have to do low-level configuration or the selection and management of hardware, an operating system, and many of the Elasticsearch and Kibana configuration parameters.

There are multiple cloud providers that provide managed clusters for Elastic Stack, such as Elastic Cloud, QBox.io, Bonsai, and many more. In this section, we will go through how to get started with **Elastic Cloud**. Elastic Cloud is the official cloud offering by the company Elastic.co, which is the main company contributing to the development of Elasticsearch and other Elastic Stack components. We will cover the following topics while working with Elastic Cloud:

- Getting up and running on Elastic Cloud
- Using Kibana
- Overriding configuration
- Recovering from a snapshot

Getting up and running on Elastic Cloud

Sign up for Elastic Cloud using <https://www.elastic.co/cloud/as-a-service/signup>, provide your email address, and verify your email. You will be asked to set your initial password.

After your initial password is set, you can log in to the Elastic Cloud console at <https://cloud.elastic.co>. The Elastic Cloud console offers an easy-to-use user interface to manage your clusters. Since you just signed up for a trial account, you can create a free cluster during the trial period.

We can choose a name for your trial cluster. You will also be able to choose **AWS** (Amazon Web Services) or **GCE** (Google Compute Engine) while launching the cluster. Upon logging in, you can create a cluster from the following screen:

pranav.shukla@gmail.com Trial started

Deployments
Custom plugins
Account
Help

Create deployment

- ### 1 Name your deployment

Give your deployment a name
- ### 2 Select a cloud platform

Pick your cloud and let us handle the rest. No additional accounts required.

Amazon Web Services

Google Cloud Platform
- ### 3 Select a region

US East (N. Virginia)

US West (N. California)

US West (Oregon)

EU (Ireland)

Asia Pacific (Singapore)

Asia Pacific (Tokyo)

South America (Sao Paulo)

Asia Pacific (Sydney)

EU (Frankfurt)
- ### 4 Set up your deployment

Stable versions

☒ 7.1.0
 ☐ 7.0.1

☐ 6.8.0

☐ 5.6.16

☐ Select a deployment to restore from its latest snapshot
- ### 5 Optimize your deployment

I/O Optimized
Recommended
 Use for search and general all-purpose workloads. Includes a balance of compute, memory, and storage.
[Default specs](#)

Compute Optimized
 Run CPU-intensive workloads or run smaller workloads cost-effectively when you need less memory and storage.
[Default specs](#)

Memory Optimized
 Perform memory-intensive operations efficiently, including workloads with frequent aggregations.
[Default specs](#)

UNAVAILABLE
Hot-Warm Architecture
 Use for time-series analytics and logging workloads that benefit from automatic index curation.
[Default specs](#)

Elastic Cloud supports many more options to cater to your specific use case such as hot-warm architecture optimized for logging, compute-focused setup optimized for analytics etc. [Learn more ...](#)

Deployment pricing

Free! As part of your 14-day trial, you can try it out without a credit card. If you want to unleash the full power of Elasticsearch Service now, you can enter your credit card details or contact sales@elastic.co.

Create deployment

Customize deployment

Fig-9.1: Creating a new cluster on Elastic Cloud

After selecting the cloud platform, you can choose a region for your cluster.

Select the version to be the latest 7.x version that is available. At the time of writing this course, version 7.1.0 is the latest version available on Elastic Cloud. You have the option of choosing either **I/O Optimized**, **Compute Optimized**, **Memory Optimized**, or **Hot-Warm Architecture** deployment. Different types of clusters are suitable for different use cases.

When you click the `Create deployment` button, your cluster will be created and started with production-grade configuration. The cluster will be secured. It will also start with a Kibana instance. At this point, it should provide you with a username/password to be used for logging into your Elasticsearch and Kibana nodes. Please note it down. It also provides a `Cloud ID`, which is a helpful string when connecting to your cloud cluster from your Beats agents and Logstash servers.

You can click under the **Deployments** text where you will see the name with which you created your deployment. In this case, we called it `test-cluster`. If you click on that, you should see a screen that has a summary of your deployment:

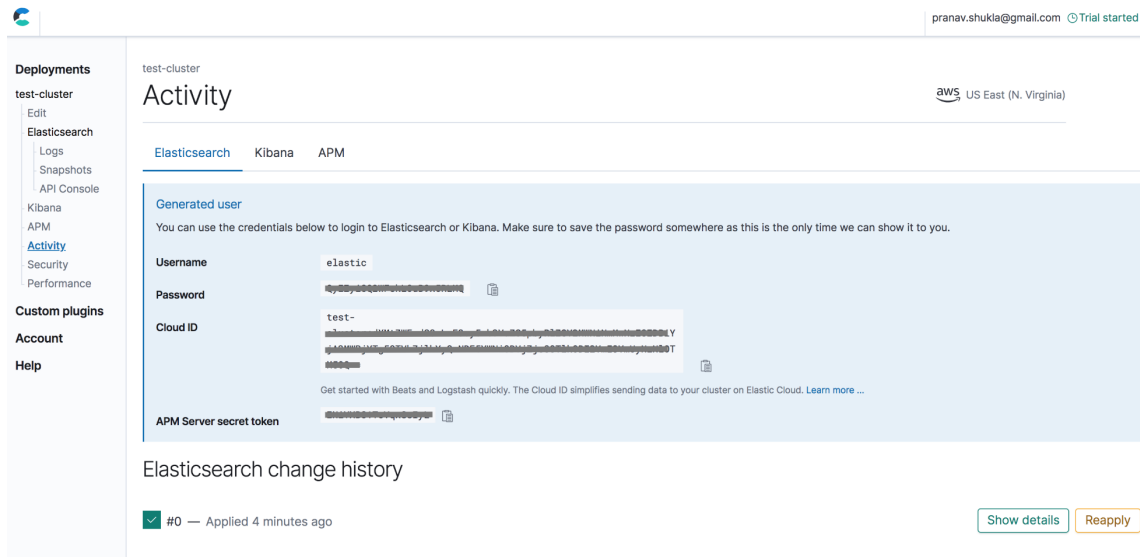


Fig-9.2: Deployment Overview screen on Elastic Cloud

As you can see, the cluster is up and running. In the second tab, **Kibana**, you can get the URL at which it is accessible. The Elasticsearch cluster is available at the given secured HTTPS URL.

The cluster has two nodes: one in each AWS availability zone and one tiebreaker node. The tiebreaker node helps to elect a master node. Tiebreaker nodes are special nodes on Elastic Cloud that help in the re-election of masters whenever some nodes become unreachable in the cluster.

Now that we have the cluster up and running with a Kibana instance, let's use it!

Using Kibana

The link to the Kibana instance is already made available to us on the cluster overview page on Elastic Cloud. You can click on it to launch the Kibana UI. Unlike the local instance of Kibana that we initially created, this instance is secured by X-Pack security. You will have to log in using the credentials provided to you after you created the Elastic Cloud cluster in the previous section.

After logging in, you should see the **Kibana** UI, as follows:

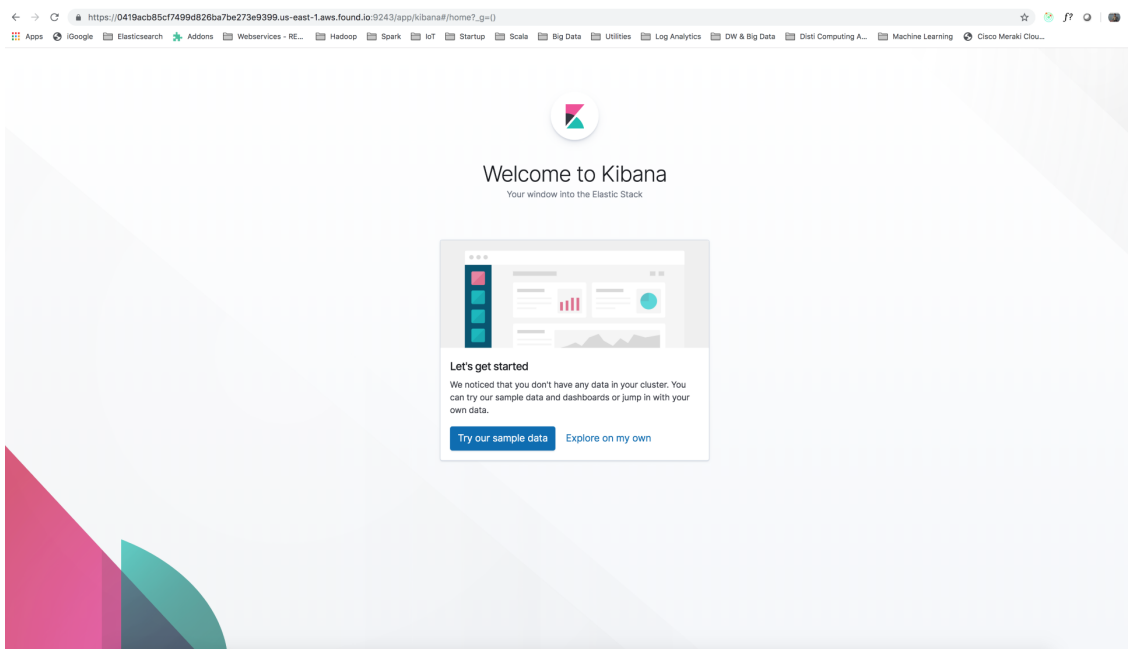


Fig-9.3 Kibana UI on Elastic Cloud after logging in

You can view all indexes, analyze data on your Elasticsearch cluster, and monitor your Elasticsearch cluster from this Kibana UI.

Overriding configuration

It is possible to override the configuration of your Elasticsearch nodes via the `Edit` menu in the navigation panel on the left side under the **Deployments**. Elastic Cloud doesn't allow you to edit the `elasticsearch.yml` file directly. However, it provides a section called **User Settings**, which allows you to override a subset of the configuration parameters.

The configuration parameters that can be overridden are documented in the Elastic Cloud reference documentation at <https://www.elastic.co/guide/en/cloud/current/ec-add-user-settings.html>.

Recovering from a snapshot

Elastic Cloud automatically creates a snapshot of all indexes in your cluster periodically (every 30 minutes) and keeps them for recovery purposes, if required. This happens automatically without doing any additional setup or code. You can visit the `Snapshots` link under your **Deployments** > **Elasticsearch** to view the available list of **Snapshots**, as follows:

Deployments

test-cluster

Edit

Elasticsearch

Logs

Snapshots

API Console

Kibana

APM

Activity

Security

Performance

Custom plugins

Account

Help

test-cluster / Elasticsearch

Snapshots

US East (N. Virginia)

Snapshots are backups of your data that you can restore in the event of an unexpected data loss.

Last successful snapshot

Next snapshot

Snapshot frequency

19 minutes ago

In 11 minutes

30 minutes

The current settings retain 100 snapshots over a period of 2 days. [Edit settings](#)

Take snapshot now

Restore from another deployment

Choose an snapshot in another deployment to restore on this deployment

Snapshots

All snapshots 29 Success 29

Snapshot	Status	Completed	Duration	
scheduled-1558940570-instance-000000001	Success	19 minutes ago	A few seconds	Restore
scheduled-1558938768-instance-000000001	Success	An hour ago	A few seconds	Restore
scheduled-1558936966-instance-000000001	Success	An hour ago	A few seconds	Restore
scheduled-1558935163-instance-000000001	Success	2 hours ago	A few seconds	Restore
scheduled-1558933361-instance-000000001	Success	2 hours ago	A few seconds	Restore
scheduled-1558931559-instance-000000001	Success	3 hours ago	A few seconds	Restore
scheduled-1558929757-instance-000000001	Success	3 hours ago	A few seconds	Restore
scheduled-1558927955-instance-000000001	Success	4 hours ago	A few seconds	Restore
scheduled-1558926153-instance-000000001	Success	4 hours ago	A few seconds	Restore
scheduled-1558924351-instance-000000001	Success	5 hours ago	A few seconds	Restore

< 1 2 3 >

Fig-9.4: Listing of snapshots on Elastic Cloud

You can choose the snapshot that you want to restore from, and you will be presented with the following screen:

The screenshot shows the 'Deployments' page for a cluster named 'test-cluster'. The left sidebar contains navigation links for Deployments, Custom plugins, Account, and Help. The main content area displays details for a specific snapshot: 'scheduled-1558940570-instance-0000000001'. It indicates the snapshot was successful, started 21 minutes ago, and completed 21 minutes ago with a duration of 755ms. Below this, it shows 4 total shards, all of which are successful, with 0 failed shards. The snapshot's UUID is 'rr0WY62QR7~apgKg1fpgXg' and its version is '7.1.0'. A list of indices is provided: '.kibana_1', '.kibana_task_manager', '.security-7', and 'apm-7.1.0-onboarding-2019.05.26'.

The 'Restore Snapshot' section is active, showing a form to specify and rename indices. The 'Specify Indices' field contains 'e.g. index1,index2'. The 'Rename Indices' section has a 'Match' field with 'index_(.+)' and a 'Replace with' field with 'restored_index_\$1'. A red 'Restore snapshot' button is at the bottom.

Fig-9.5: Snapshot details and restoring from a specific snapshot

The snapshot contains the saved state for all indexes in the cluster. It is possible to choose a subset of the indexes for restoring and also to rename it while restoring it. It is also possible to restore the snapshot on a separate cluster.

Next, we will see how to get started with Elastic Stack if you are planning to manage the Elastic Stack components yourself. This is also called **self-hosting**, in that you will be hosting and managing it on your own.

Hosting Elastic Stack on your own

Hosting Elastic Stack on your own, that is, self-hosting Elastic Stack, requires you to install, configure, and manage Elasticsearch and your other Elastic Stack products. This can be done in one of two ways:

- Self-hosting on-premise
- Self-hosting on a cloud

Regardless of whether you run Elastic Stack on-premise (in your own data center) or run it on one of the cloud providers, such as AWS, Azure, or GCE, there are some common aspects that you should take into consideration. While self-hosting, you will be faced with the following choices:

- Selecting hardware
- Selecting the operating system

- Configuring Elasticsearch nodes
- Managing and monitoring Elasticsearch nodes
- Special considerations while self-hosting on a cloud

Except for the last item, which is applicable only if you are self-hosting on a cloud, the others are equally applicable for cloud and on-premise deployments.

Selecting hardware

Elasticsearch primarily has memory-bound tasks which rely on the inverted index. The more data that it can fit in the RAM, the faster the performance will be. But this statement cannot always be generalized. It depends on the nature of your data and the type of operations or workload that you are going to have.

Using Elasticsearch doesn't mean that it has to perform all operations in-memory. Elasticsearch also uses on-disk data very efficiently, especially for aggregation operations.

Note

All datatypes (except analyzed strings) support a special data structure called `doc_values`, which organizes the data on the disk in a columnar fashion. `doc_values` is useful for sorting and aggregation operations. Since `doc_values` is enabled by default for all datatypes except analyzed strings, it makes sorts and aggregations run mostly off the disk. Those fields do not need to be loaded in memory to aggregate or sort by them.

As Elasticsearch can scale horizontally, this is a relatively easy decision to make. It is fine to start with nodes of around 16 or 32 GB RAM, with around 8 CPU cores. As we will see in the coming sections, you cannot have Elasticsearch JVM with more than 32 GB of heap; effectively, there is no point in having a machine with more than 64 GB RAM. SSD hard disks are recommended if you are planning to do heavy aggregations.

It is important to benchmark with the initial hardware and then add more nodes or upgrade your nodes.

Selecting an operating system

Linux is the preferred choice when deploying Elasticsearch and the Elastic Stack components. Your choice of operating system will mostly depend on the preferred technologies of your organization. Elastic Stack can also be deployed on Windows if your organization prefers the Microsoft stack.

Configuring Elasticsearch nodes

Elasticsearch, which is the heart of Elastic Stack, needs some configuration before starting it in production. Most of the configuration should work out of the box, but will require the following things to be reviewed at the OS level or JVM level.

JVM heap size

Set `-Xms` and `-Xmx` to be the same. More heap means Elasticsearch can keep more data in memory for faster access. But more heap also means that when the Java heap is close to full, the JVM's garbage collector will run a full garbage collection. At that point, all other processing within the Elasticsearch node experiences a pause. So, the larger the heap, the longer the pauses will be. The maximum heap size that you can configure is around 32 GB. Another recommendation to keep in mind is that we should allocate no more than 50% of the total available RAM on the machine to the Elasticsearch JVM. The reason is that the system needs enough memory for the filesystem cache for Apache Lucene. Ultimately, all the data stored on the Elasticsearch node is managed as Apache Lucene indexes, which needs RAM for fast access to the files.

So, if you are planning to store huge amounts of data in Elasticsearch, there is no point in having one single node with more than 64 GB RAM (50% of which is 32 GB, the maximum heap size). Instead, add more nodes if you want to scale.

Disable swapping

When swapping is enabled, an OS generally has a tendency to reclaim the memory from an application by swapping the data to disk to make more memory available for other programs.

On the Elasticsearch node, this can result in the OS swapping out the heap memory of Elasticsearch. This process of swapping out from memory to disk and then swapping back from disk to memory can slow down the process. This is why swapping should be disabled on the node that is running Elasticsearch.

File descriptors

On Linux and macOS operating systems, there is a limit to the number of open file handles or file descriptors that a process can keep. This often needs to be increased in the case of Elasticsearch, as the default value is generally quite low for the open file descriptor limit.

Thread pools and garbage collector

Elasticsearch does many types of operations, such as indexing, searching, sorting, and aggregations, and uses JVM thread pools to accomplish its tasks. It is advisable to not tune the settings related to thread pools in Elasticsearch. They generally do more harm than help to improve performance. Another thing not to tune in Elasticsearch is the garbage collector settings.

Managing and monitoring Elasticsearch

When you self-host Elasticsearch, the entire monitoring and management activities for the cluster are on you. It is necessary to monitor your Elasticsearch node process status, memory, and disk space on the node. If a node crashes for any reason, or becomes unavailable, it needs to be started back again.

The snapshots of the Elasticsearch indexes need to be taken regularly for taking backups. We will discuss the snapshot/restore functionalities for backing up. Most of the monitoring can be achieved via X-Pack and Kibana, but management processes need to be set up manually.

Running in Docker containers

Docker is a popular way of containerizing and shipping software. The advantage of Docker is that the software that is dockerized and runs inside a light-weight container that has a small overhead compared to a virtual machine. As a result of its reduced overhead and large pool of publicly available Docker images, Docker is a great way to run software in production in a predictable way without the need of much configuration.

Official Elasticsearch Docker images are available for download in different flavors:

- Elasticsearch with basic X-Pack license
- Elasticsearch with full X-Pack license and 30-day evaluation
- Open source version of Elasticsearch without X-Pack

Getting started with an Elasticsearch instance running inside Docker is as easy as installing Docker and running the `docker pull` command with the Elasticsearch image of your choice. The following simple commands will get your single-node Elasticsearch 7.0.1 up and running if you have Docker installed on your system:

```
docker pull docker.elastic.co/elasticsearch/elasticsearch:7.0.1

docker run -p 9200:9200 -p 9300:9300 -e "discovery.type=single-node"
docker.elastic.co/elasticsearch/elasticsearch:7.0.1
```

Docker is a highly recommended way of running applications in a predictable way in production. You can find out more about how to run Elasticsearch in Docker in a production environment in the documentation---

<https://www.elastic.co/guide/en/elasticsearch/reference/7.0/docker.html>.

Special considerations while deploying to a cloud

While self-hosting on a cloud, you may choose one of the cloud providers, such as AWS, Microsoft Azure, or GCE. They provide compute resources, networking capabilities, virtual private clouds, and much more, to get control over your servers. Using a cloud provider as opposed to running on your own hardware comes with the following advantages:

- No upfront investment in hardware
- Ability to upgrade/downgrade servers
- Ability to add or remove servers as and when needed

It is typical to not be sure how much CPU, RAM, and so on, is required for your nodes when you start. Choosing the cloud gives the flexibility to benchmark on one type of configuration and then upgrade/downgrade or add/remove nodes as needed without incurring upfront costs. We will take EC2 as an example and try to understand the considerations to take into account. Most of the considerations should remain similar for other cloud providers as well. The following are some of the aspects to consider on AWS EC2:

- Choosing instance type
- Changing the ports; do not expose ports!
- Proxy requests
- Binding HTTP to local addresses
- Installing EC2 discovery plugin
- Installing S3 repository plugin
- Setting up periodic snapshots

Let's focus on them one by one.

Choosing instance type

EC2 offers different types of instances to meet different requirements. A typical starting point for Elasticsearch is to consider the `m5d.large` or `m5d.2xlarge` instance; they have 4 CPU cores and 8 CPU cores with 16 and 32 GB RAM respectively, and SSD storage. It is always good to benchmark on your data and monitor the resource usage on your nodes. You can upgrade or downgrade the nodes as per your findings.

Changing default ports; do not expose ports!

Running any type of service in a cloud involves different security risks. It is important that none of the ports used by Elasticsearch are exposed and accessible from the public internet. EC2 allows detailed control over which ports are accessible and from which IP addresses or subnets. Generally, you should not need to make any ports accessible from outside anywhere other than port `22` in order to log in remotely.

By default, Elasticsearch uses port `9200` for HTTP traffic and `9300` for inter-node communication. It is advisable to change these default ports by editing `elasticsearch.yml` on all nodes.

Proxy requests

Use a reverse proxy such as nginx (pronounced ****engine x**) or Apache to proxy your requests to Elasticsearch/Kibana.

Binding HTTP to local addresses

You should run your Elasticsearch nodes in a **VPC** (Virtual Private Cloud). More recently, AWS creates all nodes in a VPC. The nodes that do not need to interface with the clients accept the queries from clients over HTTP. This can be done by setting `http.host` in `elasticsearch.yml`. You can find out more about the HTTP host/port bindings in the reference documentation at <https://www.elastic.co/guide/en/elasticsearch/reference/current/modules-http.html>.

Installing EC2 discovery plugin

Elasticsearch nodes discover their peers via multicast when they are in the same network. This works very well in a regular LAN. When it comes to EC2, the network is shared and the node to node communication and automatic discovery don't work. It requires the installation of the EC2 discovery plugin on all nodes to be able to discover new nodes.

To install the EC2 discovery plugin, follow the instructions at <https://www.elastic.co/guide/en/elasticsearch/plugins/current/discovery-ec2.html> and install it on all nodes.

Installing the S3 repository plugin

It is important to back up your data in Elasticsearch regularly to restore the data if a catastrophic event occurs or if you want to revert to a last known healthy state. We will look at how to backup and restore using the snapshot/restore APIs of Elasticsearch in the next section. In order to take regular backups and store them in centralized and resilient data storage, we need to set up a snapshot mechanism. When you are running Elasticsearch in EC2, it makes sense to store snapshots in an AWS S3 bucket.

Note

S3 stands for **Simple Storage Service**. It is a scalable, durable, and reliable storage service to store large amounts of data. It provides comprehensive security for your data and accessibility from many different platforms. It can meet very stringent compliance requirements due to its comprehensive security support. It is often the preferred solution for storing long-term data, especially when systems that generate the data are hosted on AWS.

The S3 repository plugin can be installed using the following command; it needs to be installed on every node of your Elasticsearch cluster:

```
sudo bin/elasticsearch-plugin install repository-s3
```

Setting up periodic snapshots

Once you have a repository set up on S3, we need to ensure that actual snapshots are taken periodically. What this means is that we need a scheduled job that triggers the command to take a snapshot at regular intervals. The interval could be 15 minutes, 30 minutes, one hour, and so on, depending on the sensitivity of your data. We will see how to establish the snapshot/restore process for your cluster in depth later in this lab.

These are some of the considerations that you have to address while running Elasticsearch in production on AWS or other clouds.

So far, we have covered how to get your production up and running on a managed cloud or self-hosted environment. If you opted to self-host, you will need to set up a backup and restore process so that you don't lose your data. The next section is only applicable if you are self-hosting your Elasticsearch cluster.

Backing up and restoring

Taking regular backups of your data to recover in the event of catastrophic failures is absolutely critical. It is important that all of your data is saved periodically at fixed time intervals and a sufficient number of such backups are preserved.

A common strategy is to take a full backup of your data at regular intervals and keep a fixed number of backups. Your cluster may be deployed on-premise in your own data center or it may be deployed on a cloud hosted service such as AWS, where you may be managing the cluster yourself.

We will look at the following topics on how to manage your backups and restore a specific backup if it is needed:

- Setting up a repository for snapshots

- Taking snapshots
- Restoring a specific snapshot

Let's look at how to do these one by one.

Setting up a repository for snapshots

The first step in setting up a regular backup process is setting up a repository for storing snapshots. There are different places where we could store snapshots:

- A shared filesystem
- Cloud or distributed filesystems (S3, Azure, GCS, or HDFS)

Depending upon where the Elasticsearch cluster is deployed, and which storage options are available, you may want to set up the repository for your snapshots in a certain way.

Let's first understand how you would do this in the simplest of scenarios, when you want to store it in a shared filesystem directory.

Shared filesystem

When your cluster has a shared filesystem accessible from all the nodes of the cluster, you have to ensure that the shared filesystem is accessible on a common path. You should mount that shared folder on all nodes and add the path of the mounted directory. The shared, mounted filesystem's path should be added to each node's

`elasticsearch.yml` as follows:

```
path.repo: ["/mount/es_backups"]
```

Note

If you are running a single node cluster and haven't set up a real distributed cluster, there is no need for a mounted shared drive. The `path.repo` parameter can be set to a local directory of your node. It is not recommended to run a production server on a single node cluster.

Once this setting is added to `config/elasticsearch.yml` on all nodes, please restart all the nodes of your cluster.

The next step is to register a named repository under this registered folder. This is done using the following `curl` command, where we are registering a named repository with the name `backups` :

```
curl -XPUT 'http://localhost:9200/_snapshot/backups' -H 'Content-Type: application/json' -d '{
  "type": "fs",
  "settings": {
    "location": "/mount/es_backups/backups",
    "compress": true
  }
}'
```

You will need to replace `localhost` with the hostname or IP address of one of the nodes on your cluster.

The `type` parameter set to `fs` is for the shared filesystem. The `settings` parameter's body depends on the `type` parameter's value.

Since we are currently looking at a shared filesystem snapshot repository, the body of the `settings` parameter has specific parameters to set up the shared filesystem-based repository. If the `location` parameter is specified as an absolute path, it must be under one of the folders registered with the `path.repo` parameter

in `elasticsearch.yml`. If the `location` parameter is not an absolute path, Elasticsearch will assume it is a relative path from the `path.repo` parameter. The `compress` parameter saves the snapshots in compressed format.

Cloud or distributed filesystems

When you are running your Elasticsearch cluster on AWS, Azure, or Google Cloud, it makes sense to store the snapshots in one of the alternatives provided by the cloud platform to store the data in robust, fault tolerant storage, rather than storing it on a shared drive.

Elasticsearch has official plugins that allow you to store the snapshots in S3. All you need to do is install the repository--- `s3` plugin on all nodes of your cluster and set up the repository settings in a similar way to how we set up the shared filesystem repository:

```
curl -XPUT 'http://localhost:9200/_snapshot/backups' -H 'Content-Type: application/json' -d '{
  "type": "s3",
  "settings": {
    "bucket": "bucket_name",
    "region": "us-west",
    ...
  }
}'
```

The `type` should be `s3` and `settings` should have relevant values for `s3`.

Taking snapshots

Once the repository is set up, we can put named snapshots into a specific repository:

```
curl -XPUT 'http://localhost:9200/_snapshot/backups/backup_201905271530?pretty' -H 'Content-Type: application/json' -d'
{
  "indices": "bigginsight,logstash-*",
  "ignore_unavailable": true,
  "include_global_state": false
}
'
```

In this command, we specified that we want a snapshot to be taken in the repository `backups` with the name `backup_201905271530`. The name of the snapshot could be anything, but it should help you identify the snapshot at a later stage. One typical strategy would be to take a snapshot every 30 minutes and set snapshot names with prefixes such as `backup_yyyyMMddHHmm`. In the event of any failure, you could then identify the snapshot that can be restored.

Snapshots are incremental by default. They don't store all the redundant data in all snapshots.

Having taken the snapshots periodically, you would want to list all the snapshots that exist in a repository. This can be done using the following command:

```
curl -XGET 'http://localhost:9200/_snapshot/backups/_all?pretty'
```

Restoring a specific snapshot

If the need arises, you can restore the state from a specific snapshot using the following command:

```
curl -XPOST 'http://localhost:9200/_snapshot/backups/backup_201905271530/_restore'
```

This will restore the `backup_201905271530` snapshot from the `backups` repository.

Once we have set up a periodic job that takes and stores a snapshot, we are safe in the event of any failure. We now have a cluster that is recoverable from any disaster-like situation. Remember, the output of snapshots should be stored in resilient storage. At least, it should not be saved on the same Elasticsearch cluster; it should be saved on different storage, preferably a robust filesystem that is highly available, such as S3, HDFS, and so on.

So far in this lab, we have got up and running with a cluster that is reliable and is backed up regularly. In the upcoming sections, we will see how to address some common scenarios in data modeling. We will see some common strategies for setting up aliases for indexes, index templates, modeling time-series data, and so on.

Setting up index aliases

Index aliases let you create aliases for one or more indexes or index name patterns. We will cover the following topics in order to learn how index aliases work:

- Understanding index aliases
- How index aliases can help

Understanding index aliases

An index alias just provides an extra name to refer to an index; it can be defined in the following way:

```
POST /_aliases
{
  "actions" : [
    { "add" : { "index" : "index1", "alias" : "current_index" } }
  ]
}
```

Here, `index1` can be referred to with the alias `current_index`. Similarly, the index alias can be removed with the `remove` action of the `_aliases` REST API:

```
POST /_aliases
{
  "actions" : [
    { "remove" : { "index" : "index1", "alias" : "current_index" } }
  ]
}
```

The preceding call will remove the alias `current_index`. Two actions can be combined in a single invocation of the `_aliases` API. When two calls are combined, the operations are done automatically. For example, the following call would be completely transparent to the client:

```
POST /_aliases
{
  "actions" : [
    { "remove" : { "index" : "index1", "alias" : "current_index" } },
    { "add" : { "index" : "index2", "alias" : "current_index" } }
  ]
}
```

```
]
}
```

Before the call, the alias `current_index` was referring to the index `index1` `[*,*]` and after the call, the alias will refer to the index `index2`.

How index aliases can help

Once in production, it often happens that we need to reindex data from one index to another. We might have one or more applications developed in Java, Python, .NET, or other programming environments that may be referring to these indexes. In the event that the production index needs to be changed from `index1` to `index2`, it will require a change in all client applications.

Aliases come to the rescue here. They offer extra flexibility, and hence, they are a recommended feature to use in production. The key thing is to create an alias for your production index and use the alias instead of the actual index name in the client applications that use them.

In the event that the current production index needs to change, we just need to update the alias to point to the new index instead of the old one. Using this feature, we can achieve zero downtime in production in the case of data migration or the need for reindexing. Aliases use a famous principle in computer science---an extra layer of indirection can solve most problems in computer science---<https://en.wikipedia.org/wiki/Indirection>.

Apart from the ones discussed here, there are more features that aliases offer; these include the ability to use index patterns, routing, the ability to specify filters, and many more. We will see how index aliases can be leveraged when creating time-based indexes later in the lab.

Setting up index templates

One important step while setting up your index is defining the mapping for the types, number of shards, replica, and other configurations. Depending upon the complexity of the types within your index, this step can involve a substantial amount of configuration.

Index templates allow you to create indexes based on a given template, rather than creating each index manually beforehand. Index templates allow you to specify settings and mappings for the index to be created. Let's understand this by going through the following points:

- Defining an index template
- Creating indexes on the fly

Let's say we want to store sensor data from various devices and we want to create one index per day. At the beginning of every day, we want a new index to be created whenever the first sensor reading is indexed for that day. We will look into the details of why we should use such time-based indexes in the next section.

Defining an index template

We start by defining an index template:

```
PUT _template/readings_template      1
{
  "index_patterns": ["readings*"],    2
  "settings": {                       3
    "number_of_shards": 1
  },
  "mappings": {                       4
    "properties": {
```

```

    "sensorId": {
      "type": "keyword"
    },
    "timestamp": {
      "type": "date"
    },
    "reading": {
      "type": "double"
    }
  }
}
}

```

In this `_template` call, we define the following things:

- A template with the name `readings_template`.
- The index name patterns that will match this template. We configured `readings*` as the one and only index pattern. Any attempt to index into an index that does not exist but matches this pattern would use this template.
- The settings to be applied to the newly created index from this template.
- The mappings to be applied to the newly created index from this template.

Let's try to index data into this new index.

Creating indexes on the fly

When any client tries to index the data for a particular sensor device, it should use the index name with the current day appended in `yyyy-mm-dd` format after `readings`. A call to index data for `2019-05-01` would look like the following:

```

POST /readings-2019-05-01/_doc
{
  "sensorId": "a11111",
  "timestamp": 1483228800000,
  "reading": 1.02
}

```

When the first record for the date `2019-05-01` is being inserted, the client should use the index name `readings-2019-05-01`. Since this index doesn't exist yet, and we have an index template in place, Elasticsearch creates a new index using the index template we defined. As a result, the settings and mappings defined in our index template get applied to this new index.

This is how we create indexes based on index templates. In the next section, let's understand why these types of time-based indexes are useful and how to use them in production with your time-series data.

Modeling time series data

Often, we have a need to store time series data in Elasticsearch. Typically, one would create a single index to hold all documents. This typical approach of one big index to hold all documents has its own limitations, especially for the following reasons:

- Scaling the index with an unpredictable volume over time
- Changing the mapping over time
- Automatically deleting older documents

Let's look at how each problem manifests itself when we choose a single monolithic index.

Scaling the index with unpredictable volume over time

One of the most difficult choices when creating an Elasticsearch cluster and its indexes is deciding how many primary shards should be created and how many replica shards should be created.

Let's understand how the number of shards becomes important in the following sub-sections:

- Unit of parallelism in Elasticsearch:
 - The effect of the number of shards on the relevance score
 - The effect of the number of shards on the accuracy of aggregations

Unit of parallelism in Elasticsearch

We have to decide the number of shards at the time of creating the index. The number of shards cannot be changed once the index has been created. There is no golden rule that will help you decide how many shards should be created at the time of creating an index. The number of shards actually decides the level of parallelism in the index. Let's understand this by taking an example of how a search query might be executed.

When a search or aggregation query is sent by a client, it is first received by one of the nodes in the cluster. That node acts as a coordinator for that request. The coordinating node sends requests to all the shards on the cluster and waits for the response from all shards. Once the response is received by the coordinating node from all shards, it collates the response and sends it back to the original client.

What this means is, when we have a greater number of shards, each shard has to do relatively less work and parallelism can be increased.

But can we choose an arbitrarily big number of shards? Let's look at this in the next couple of sub-sections.

The effect of the number of shards on the relevance score

A large number of small shards is not always the solution, as it can affect the relevance of the search results. In the context of search queries, the relevance score is calculated within the context of a shard. The relative frequencies of documents are calculated within the context of each shard and not across all shards. This is why the number of shards can affect the overall scores observed for a query. In particular, having too many shards to address the future scalability problem is not a solution.

The effect of the number of shards on the accuracy of aggregations

Similar to the execution of the search query, an aggregation query is also coordinated by a coordinating node. Let's say that the client has requested terms aggregation on a field that can take a large number of unique values. By default, the terms aggregation returns the top 10 terms to the client.

To coordinate the execution of terms aggregation, the coordinator node does not request all the buckets from all shards. All shards are requested to give their top $[n]$ buckets. By default, this number, $[n]$, is equal to the `size` parameter of the terms aggregation, that is, the number of top buckets that the client has requested. So, if the client requested the top 10 terms, the coordinating node in turn requests the top 10 buckets from each shard.

Since the data can be skewed across the shards to a certain extent, some of the shards may not even have certain buckets, even though those buckets might be one of the top buckets in some shards. If a particular bucket is in the top $[n]$ buckets returned by one of the shards and that bucket is not one of the top $[n]$ buckets by one of the other shards, the final count aggregated by the coordinating node will be off for that bucket. A large number of shards, just to ensure future scalability, does not help the accuracy of aggregations.

We have understood why the number of shards is important and how deciding the number of shards upfront is difficult. Next, we will see how changing the mapping of indexes becomes difficult over a period of time.

Changing the mapping over time

Once an index is created and documents start getting stored, the requirements can change. There is only one thing that is constant, **change**.

When the schema changes, the following types of change may happen with respect to the schema:

- New fields get added
- Existing fields get removed

New fields get added

When the first document with a new field gets indexed, the new field's mapping is automatically created if it doesn't already exist. Elasticsearch infers the datatype of the field based on the value of that field in the first document in order to create the mapping. The mappings of one particular type of document can grow over a period of time.

Once a document with a new field is indexed, the mapping is created for that new field and its mapping remains.

Existing fields get removed

Over a period of time, the requirements of a project can change. Some fields might become obsolete and may no longer be used. In the case of Elasticsearch indexes, the fields that are no longer used are not removed automatically; the mapping remains in the index for all the fields that were ever indexed. Each extra field in the Elasticsearch index carries an overhead; this is especially true if you have hundreds or thousands of fields. If, in your use case, you have a very high number of fields that are not used, it can increase the burden on your cluster.

Automatically deleting older documents

No cluster has an infinite capacity to retain data forever. With the volume growing over a period of time, you may decide to only store necessary data in Elasticsearch. Typically, you may want to retain data for the past few weeks, months, or years in Elasticsearch, depending on your use case.

Prior to Elasticsearch 2.x, this was achieved using **TTL** (Time to Live) set on individual documents. Each document could be configured to remain in the index for a configurable amount of time. But, the TTL feature was deprecated with the 2.x version because of its overheads in maintaining time-to-live on a per-document basis.

We have seen some problems that we might face while dealing with time series data. Now, let's look at how the use of **time-based indexes** addresses these issues. Time-based indexes are also called **index-per-timeframe**:

- How index-per-timeframe solves these issues
- How to set up index-per-timeframe

How index-per-timeframe solves these issues

Instead of going with one big monolithic index, we now create one index per timeframe. The timeframe could be one day, one week, one month, or any arbitrary time duration. For example, in our example in the *[Index Template]* section, we chose index-per-day. The names of the index would reflect that--we had indexes such as `readings-2019-05-01`, `readings-2019-05-02`, and so on. If we had chosen index-per-month, the index names would look like `readings-2019-04`, `readings-2019-05`, `readings-2019-06`, and so on.

Let's look at how this scheme solves the issues we saw earlier one by one.

Scaling with index-per-timeframe

Since we no longer have a monolithic index that needs to hold all historic data, scaling up or scaling down according to the recent volumes becomes easier. The choice of the number of shards is not an upfront and permanent decision. Start with an initial estimated number of shards for the given time period. This number, the chosen number of shards, can be put in the index template.

Since that choice of shards can be changed before the next timeframe begins, you are not stuck with a bad choice. With each time period, it gives a chance to adjust the index template to increase or decrease the number of shards for the next index to be created.

Changing the mapping over time

Changing the mapping becomes easier, as we could just update the index template that is used for creating new indexes. When the index template is updated, the new index that is created for the new timeframe uses the new mappings in the template.

Again, each timeframe gives us an opportunity to change.

Automatically deleting older documents

With time-based indexes, deleting the older documents becomes easier. We could just drop older indexes rather than delete individual documents. If we were using monthly indexes and wanted to enforce six-month retention of data, we could delete all indexes older than 6 months. This may be set up as a scheduled job to look for and delete older indexes.

As we have seen in this section, setting up index-per-timeframe has obvious advantages when we are dealing with time-series data.

Summary

In this lab, we have seen essential techniques necessary to take your next Elastic Stack application to production. We have seen various deployment options, including cloud-based and on-premise. We have seen how to use a managed cloud service provider such as Elastic Cloud and have also covered how to self-host Elastic Stack. We have covered some common concerns and decision choices that you will face, whether you self-host or use a managed cloud provider.

Additionally, we have seen various techniques useful in a production-grade Elastic Stack deployment. These include the usage of index aliases, index templates, and modeling time-series data. This is definitely not a comprehensive guide covering all the nuances of running Elastic Stack in production, but we have definitely covered enough for you to comfortably take your next Elastic Stack project to production.

Equipped with all these techniques, we will build a sensor data analytics application in the next lab, Lab 10[*Building a Sensor Data Analytics Application*] .