

Lab 10. Building a Sensor Data Analytics Application



In the previous lab, we saw how you can take an Elastic Stack application to production. Armed with all the knowledge of the Elastic Stack and the techniques for taking applications to production, we are ready to apply these concepts in a real-world application. In this lab, we will build one such application using the Elastic Stack that can handle a large amount of data, applying the techniques that we have learned so far.

We will cover the following topics as we build a sensor-data analytics application:

- Introduction to the application
- Modeling data in Elasticsearch
- Setting up the metadata database
- Building the Logstash data pipeline
- Sending data to Logstash over HTTP
- Visualizing the data in Kibana

Let's go through the topics.

Understanding the sensor-generated data

What does the data look like when it is generated by the sensor? The sensor sends JSON-format data over the internet and each reading looks like the following:

```
{
  "sensor_id": 1,
  "time": 1511935948000,
  "value": 21.89
}
```

Here, we can see the following:

- The `sensor_id` field is the unique identifier of the sensor that has emitted the record.
- The `time` field is the time of the reading in milliseconds since the epoch, that is, 00:00:00 on January 1, 1970.
- The `value` field is the actual metric value emitted by the sensor.

This type of JSON payload is generated every minute by all the sensors in the system. Since all sensors are registered in the system, the server-side system has the associated metadata with each sensor. Let's look at the sensor-related metadata that is available to us on the server side in a database.

Understanding the sensor metadata

The metadata about all the sensors across all locations is available to us in a relational database. In our example, we have stored it in MySQL. This type of metadata can be stored in any relational database other than MySQL. It can also be stored in Elasticsearch in an index.

The metadata about sensors primarily contains the following details:

- **Type of sensor:** What type of sensor is it? It can be a temperature sensor, a humidity sensor, and so on.
- **Location-related metadata:** Where is the sensor with the given sensor ID physically located? Which customer is it associated with?

This information is stored in the following three tables in MySQL:

- `sensor_type`: Defines various sensor types and their `sensor_type_id`:

sensor_type_id	sensor_type
1	Temperature
2	Humidity

- `location` : This defines locations with their latitude/longitude and address within a physical building:

location_id	customer	department	building_name	room	floor	location_on_floor	latitude	longitude
1	Abc Labs	R & D	222 Broadway	101	1	C-101	710936	-74.008500

- `sensors` : This maps `sensor_id` with sensor types and locations:

sensor_id	sensor_type_id	location_id
1	1	1
2	2	1

Given this database design, it is possible to look up all of the metadata associated with the given `sensor_id` using the following SQL query:

```
select
    st.sensor_type as sensorType,
    l.customer as customer,
    l.department as department,
    l.building_name as buildingName,
    l.room as room,
    l.floor as floor,
    l.location_on_floor as locationOnFloor,
    l.latitude,
    l.longitude
from
    sensors s
    inner join
    sensor_type st ON s.sensor_type_id = st.sensor_type_id
    inner join
    location l ON s.location_id = l.location_id
where
    s.sensor_id = 1;
```

The result of the previous query will look like this:

sensorType	customer	department	buildingName	room	floor	locationOnFloor	latitude	longitude
Temperature	Abc Labs	R & D	222 Broadway	101	Floor1	C-101	710936	-74.0085

Up until now, we have seen the format of incoming sensor data from the client side. We have also established a mechanism to look up the associated metadata for the given sensor.

Next, we will see what the final enriched record should look like.

Understanding the final stored data

By combining the data that is coming from the client side and contains the sensor's metric value for a given metric at a given time, we can construct an enriched record of the following fields:

- `sensorId`
- `sensorType`
- `customer`
- `department`
- `buildingName`
- `room`
- `floor`
- `locationOnFloor`
- `latitude`
- `longitude`
- `time`
- `reading`

Field numbers 1, 11, and 12 are present in the payload sent by the sensor to our application. The remaining fields are looked up or enriched using the SQL query that we saw in the previous section -- using the `sensorId`. This way, we can generate a denormalized sensor reading record for every sensor for every minute.

We have understood what the application is about and what the data represents. As we start developing the application, we will start the solution from the inside out. It is better to attack the problem at hand at the very heart and try to piece together its core. Elasticsearch is at the core of the Elastic Stack, so we will start defining our solution from it's very heart by first building the data model in Elasticsearch. Let's do that in the next section.

Modeling data in Elasticsearch

We have seen the structure of the final record after enriching the data. That should help us model the data in Elasticsearch. Given that our data is time series data, we can apply some of the techniques mentioned in Lab 9, *[Running the Elastic Stack in Production]*, to model the data:

- Defining an index template
- Understanding the mapping

Let's look at the index template that we will define.

Defining an index template

Since we are going to be storing time series data that is immutable, we do not want to create one big monolithic index. We'll use the techniques discussed in the *[Modeling time series data]* section in Lab 9, *[Running the Elastic Stack in Production]*.

The source code of the application in this lab is within the GitHub repository at <https://github.com/fenago/elasticsearch/tree/v7.0/lab-10>. As we go through the lab, we will perform the steps mentioned in the `README.md` file located at that path.

Please create the index template mentioned in *[Step 1]* of the `README.md` file or execute the following script in your Kibana Dev Tools Console:

```
POST _template/sensor_data_template
{
  "index_patterns": [
    "sensor_data*"
  ],
  "settings": {
    "number_of_replicas": "1",
    "number_of_shards": "5"
  },
  "mappings": {
    "properties": {
      "sensorId": {
        "type": "integer"
      },
      "sensorType": {
        "type": "keyword",
        "fields": {
          "analyzed": {
            "type": "text"
          }
        }
      },
      "customer": {
        "type": "keyword",
        "fields": {
          "analyzed": {
            "type": "text"
          }
        }
      },
      "department": {
        "type": "keyword",
        "fields": {
          "analyzed": {
            "type": "text"
          }
        }
      },
      "buildingName": {
        "type": "keyword",
        "fields": {
          "analyzed": {
```

```

        "type": "text"
      }
    }
  },
  "room": {
    "type": "keyword",
    "fields": {
      "analyzed": {
        "type": "text"
      }
    }
  },
  "floor": {
    "type": "keyword",
    "fields": {
      "analyzed": {
        "type": "text"
      }
    }
  },
  "locationOnFloor": {
    "type": "keyword",
    "fields": {
      "analyzed": {
        "type": "text"
      }
    }
  },
  "location": {
    "type": "geo_point"
  },
  "time": {
    "type": "date"
  },
  "reading": {
    "type": "double"
  }
}
}
}

```

This index template will create a new index with the name `sensor_data-YYYY.MM.dd` when any client attempts to index the first record in this index. We will see later in this lab how this can be done from Logstash under *[Building the Logstash data pipeline]* section.

Understanding the mapping

The mapping that we defined in the index template contains all the fields that will be present in the denormalized record after lookup. A few things to notice in the index template mapping are as follows:

- All the fields that contain a `text` type of data are stored as the `keyword` type; additionally, they are stored as `text` in an analyzed field. For example, please have a look at the `customer` field.

- The latitude and longitude fields that we had in the enriched data are now mapped to a `geo_point` type of field with the field name of `location`.

At this point, we have defined an index template that will trigger the creation of an index with the mapping we defined in the template.

Setting up the metadata database

We need to have a database that has metadata about the sensors. This database will hold the tables that we discussed in the *[Introduction to the application]* section.

We are storing the data in a relational database MySQL, but you can use any other relational database equally well. Since we are using MySQL, we will be using the MySQL JDBC driver to connect to the database. Please ensure that you have the following things set up on your system:

1. MySQL database community version 5.5, 5.6, or 5.7. You can use an existing database if you already have it on your system.
2. Install the downloaded MySQL database and log in with the root user. Execute the script available https://github.com/fenago/elasticsearch/tree/v7.0/lab-10/files/create_sensor_metadata.sql.
3. Log in to the newly created `sensor_metadata` database and verify that the three tables, `sensor_type`, `locations`, and `sensors`, exist in the database.

You can verify that the database was created and populated successfully by executing the following query:

```
select
  st.sensor_type as sensorType,
  l.customer as customer,
  l.department as department,
  l.building_name as buildingName,
  l.room as room,
  l.floor as floor,
  l.location_on_floor as locationOnFloor,
  l.latitude,
  l.longitude
from
  sensors s
  inner join
  sensor_type st ON s.sensor_type_id = st.sensor_type_id
  inner join
  location l ON s.location_id = l.location_id
where
  s.sensor_id = 1;
```

The result of the previous query will look like this:

sensorType	customer	department	buildingName	room	floor	locationOnFloor	latitude	longitude
Temperature	Abc Labs	R & D	222 Broadway	101	Floor1	C-101	710936	-74.0085

Our `sensor_metadata` database is ready to look up the necessary sensor metadata. In the next section, we will build the Logstash data pipeline.

Building the Logstash data pipeline

Having set up the mechanism to automatically create the Elasticsearch index and the metadata database, we can now focus on building the data pipeline using Logstash. What should our data pipeline do? It should perform the following steps:

- Accept JSON requests over the web (over HTTP).
- Enrich the JSON with the metadata we have in the MySQL database.
- Store the resulting documents in Elasticsearch.

These three main functions that we want to perform correspond exactly with the Logstash data pipeline's input, filter, and output plugins, respectively. The full Logstash configuration file for this data pipeline is in the code base at https://github.com/fenago/elasticsearch/tree/v7.0/lab-10/files/logstash_sensor_data_http.conf.

Let us look at how to achieve the end goal of our data pipeline by following the aforementioned steps. We will start with accepting JSON requests over the web (over HTTP).

Accepting JSON requests over the web

This function is achieved by the input plugin. Logstash has support for the `http` input plugin, which does precisely that. It builds an HTTP interface using different types of payloads that can be submitted to Logstash as an input.

The relevant part from `logstash_sensor_data_http.conf`, which has the input filter, is as follows:

```
input {
  http {
    id => "sensor_data_http_input"
  }
}
```

Here, the `id` field is a string that can uniquely identify this input filter later in the file if needed. We will not need to reference this name in the file; we just choose the name `sensor_data_http_input`.

The reference documentation of the HTTP input plugin is available at: <https://www.elastic.co/guide/en/logstash/current/plugins-inputs-http.html>. In this instance, since we are using the default configuration of the `http` input plugin, we have just specified `id`. We should secure this HTTP endpoint as it will be exposed over the internet to allow sensors to send data from anywhere. We can configure `user` and `password` parameters to protect this endpoint with the desired username and password, as follows:

```
input {
  http {
    id => "sensor_data_http_input"
    user => "sensor_data"
    password => "sensor_data"
  }
}
```

When Logstash is started with this input plugin, it starts an HTTP server on port `8080`, which is secured using basic authentication with the given username and password. We can send a request to this Logstash pipeline using a `curl` command, as follows:

```
curl -XPOST -u sensor_data:sensor_data --header "Content-Type: application/json"
"http://localhost:8080/" -d '{"sensor_id":1,"time":1512102540000,"reading":16.24}'
```

Let's see how we will enrich the JSON payload with the metadata we have in MySQL.

Enriching the JSON with the metadata we have in the MySQL database

The enrichment and other processing parts of the data pipeline can be done using filter plugins. We have built a relational database that contains the tables and necessary lookup data for enriching the incoming JSON requests.

Logstash has a `jdbc_streaming` filter plugin that can be used to do lookups from any relational database and enrich the incoming JSON documents. Let's zoom into the filter plugin section in our Logstash configuration file:

```
filter {
  jdbc_streaming {
    jdbc_driver_library => "/path/to/mysql-connector-java-5.1.45-bin.jar"
    jdbc_driver_class => "com.mysql.jdbc.Driver"
    jdbc_connection_string => "jdbc:mysql://localhost:3306/sensor_metadata"
    jdbc_user => "root"
    jdbc_password => "<password>"
    statement => "select st.sensor_type as sensorType, l.customer as customer,
l.department as department, l.building_name as buildingName, l.room as room, l.floor
as floor, l.location_on_floor as locationOnFloor, l.latitude, l.longitude from sensors
s inner join sensor_type st on s.sensor_type_id=st.sensor_type_id inner join location
l on s.location_id=l.location_id where s.sensor_id= :sensor_idenfier"
    parameters => { "sensor_idenfier" => "sensor_id"}
    target => lookupResult
  }

  mutate {
    rename => {"[lookupResult][0][sensorType]" => "sensorType"}
    rename => {"[lookupResult][0][customer]" => "customer"}
    rename => {"[lookupResult][0][department]" => "department"}
    rename => {"[lookupResult][0][buildingName]" => "buildingName"}
    rename => {"[lookupResult][0][room]" => "room"}
    rename => {"[lookupResult][0][floor]" => "floor"}
    rename => {"[lookupResult][0][locationOnFloor]" => "locationOnFloor"}
    add_field => {
      "location" => "%{[lookupResult][0][latitude]},%{[lookupResult][0][longitude]}"
    }
    remove_field => ["lookupResult", "headers", "host"]
  }
}
```

As you will notice, there are two filter plugins used in the file:

- `jdbc_streaming`
- `mutate`

Let's see what each filter plugin is doing.

The `jdbc_streaming` plugin

We essentially specify the whereabouts of the database that we want to connect to, the username/password, the JDBC driver `.jar` file, and the class. We already created the database in the [\[Setting up the metadata database\]](#) section.

Download the latest MySQL JDBC Driver, also known as **Connector/J**, from <https://dev.mysql.com/downloads/connector/j/>. At the time of writing this course, the latest version is 5.1.45, which works with MySQL 5.5, 5.6, and 5.7. Download the `.tar` / `.zip` file containing the driver and extract it into your system. The path of this extracted `.jar` file should be updated in the `jdbc_driver_library` parameter.

To summarize, you should review and update the following parameters in the Logstash configuration to point to your database and driver `.jar` file:

- `jdbc_connection_string`
- `jdbc_password`
- `jdbc_driver_library`

The `statement` parameter has the same SQL query that we saw earlier. It looks up the metadata for the given `sensor_id`. A successful query will fetch all additional fields for that `sensor_id`. The result of the lookup query is stored in a new field, `lookupResult`, as specified by the `target` parameter.

The resulting document, up to this point, should look like this:

```
{
  "sensor_id": 1,
  "time": 1512113760000,
  "reading": 16.24,
  "lookupResult": [
    {
      "buildingName": "222 Broadway",
      "sensorType": "Temperature",
      "latitude": 40.710936,
      "locationOnFloor": "Desk 102",
      "department": "Engineering",
      "floor": "Floor 1",
      "room": "101",
      "customer": "Linkedin",
      "longitude": -74.0085
    }
  ],
  "@timestamp": "2019-05-26T05:23:22.618Z",
  "@version": "1",
  "host": "0:0:0:0:0:0:1",
  "headers": {
    "remote_user": "sensor_data",
    "http_accept": "*/*",
    ...
  }
}
```

As you can see, the `jdbc_streaming` filter plugin added some fields apart from the `lookupResult` field. These fields were added by Logstash and the `headers` field was added by the HTTP input plugin.

In the next section, we will use the `mutate` filter plugin to modify this JSON to the desired end result that we want in Elasticsearch.

The mutate plugin

As we have seen in the previous section, the output of the `jdbc_streaming` filter plugin has some undesired aspects. Our JSON payload needs the following modifications:

- Move the looked-up fields that are under `lookupResult` directly into the JSON file.
- Combine the latitude and longitude fields under `lookupResult` as a location field.
- Remove the unnecessary fields.

```
mutate {
  rename => {"[lookupResult][0][sensorType]" => "sensorType"}
  rename => {"[lookupResult][0][customer]" => "customer"}
  rename => {"[lookupResult][0][department]" => "department"}
  rename => {"[lookupResult][0][buildingName]" => "buildingName"}
  rename => {"[lookupResult][0][room]" => "room"}
  rename => {"[lookupResult][0][floor]" => "floor"}
  rename => {"[lookupResult][0][locationOnFloor]" => "locationOnFloor"}
  add_field => {
    "location" => "%{lookupResult[0]latitude},%{lookupResult[0]longitude}"
  }
  remove_field => ["lookupResult", "headers", "host"]
}
```

Let's see how the `mutate` filter plugin achieves these objectives.

Moving the looked-up fields that are under lookupResult directly in JSON

As we have seen, `lookupResult` is an array with just one element: the element at index `0` in the array. We need to move all the fields under this array element directly under the JSON payload. This is done field by field using the `rename` operation.

For example, the following operation renames the existing `sensorType` field directly under the JSON payload:

```
rename => {"[lookupResult][0][sensorType]" => "sensorType"}
```

We do this for all the looked-up fields that are returned by the SQL query.

Combining the latitude and longitude fields under lookupResult as a location field

Remember when we defined the index template mapping for our index? We defined the `location` field to be of `geo_point` type. The `geo_point` type accepts a value that is formatted as a string with latitude and longitude appended together, separated by a comma.

This is achieved by using the `add_field` operation to construct the `location` field, as follows:

```
add_field => {
  "location" => "%{[lookupResult][0][latitude]},%{[lookupResult][0][longitude]}"
}
```

By now, we should have a new field called `location` added to our JSON payload, exactly as desired. Next, we will remove the undesirable fields.

Removing the unnecessary fields

After moving all the elements from the `lookupResult` field directly in the JSON, we don't need that field anymore. Similarly, we don't want to store the `headers` or the `host` fields in the Elasticsearch index, so we remove them all at once using the following operation:

```
remove_field => ["lookupResult", "headers", "host"]
```

We finally have the JSON payload in the structure that we want in the Elasticsearch index. Next, let us see how to send it to Elasticsearch.

Store the resulting documents in Elasticsearch

We use the Elasticsearch output plugin that comes with Logstash to send data to Elasticsearch. The usage is very simple; we just need to have `elasticsearch` under the output tag, as follows:

```
output {
  elasticsearch {
    hosts => ["localhost:9200"]
    index => "sensor_data-%{+YYYY.MM.dd}"
  }
}
```

We have specified `hosts` and `index` to send the data to the right index within the right cluster. Notice that the index name has `%{YYYY.MM.dd}`. This calculates the index name to be used by using the event's current time and formats the time in this format.

Remember that we had defined an index template with the index pattern `sensor_data*`. When the first event is sent on May 26, 2019, the output plugin defined here will send the event to index `sensor_data-2019.05.26`.

If you want to send events to a secured Elasticsearch cluster as we did when we used X-Pack in Lab 8, *[Elastic X-Pack]*, you can configure the `user` and `password` parameters as follows:

```
output {
  elasticsearch {
    hosts => ["localhost:9200"]
    index => "sensor_data-%{+YYYY.MM.dd}"
    user => "elastic"
    password => "elastic"
  }
}
```

This way, we will have one index for every day, where each day's data will be stored within its index. We had learned the index per time frame in Lab 9, *Running the Elastic Stack in Production*.

Now that we have our Logstash data pipeline ready, let's send some data.

Sending data to Logstash over HTTP

At this point, sensors can start sending their readings to the Logstash data pipeline that we have created in the previous section. They just need to send the data as follows:

```
curl -XPOST -u sensor_data:sensor_data --header "Content-Type: application/json"
"http://localhost:8080/" -d '{"sensor_id":1,"time":1512102540000,"reading":16.24}'
```

Since we don't have real sensors, we will simulate the data by sending these types of requests. The simulated data and script that send this data are incorporated in the code at <https://github.com/fenago/elasticsearch/tree/master/lab-10/data>.

If you are on Linux or macOS, open the Terminal and change the directory to your Learning Elasticsearch workspace that was checked out from GitHub.

Note

If your machine has a Windows operating system, you will need a Linux-like shell that supports the `curl` command and basic **BASH** (Bourne Again SHell) commands. As you may already have a GitHub workspace checked out, you may be using [Git for Windows, *] which has Git BASH. This can be used to run the script that loads data. If you don't have Git BASH, please download and install [Git for Windows] from <https://git-scm.com/download/win> and launch Git BASH to run the commands mentioned in this section.

Now, go to the `lab-10/data` directory and execute `load_sensor_data.sh`:

```
$ pwd
~/Desktop/elasticsearch
$ cd lab-10/data
$ ls
load_sensor_data.sh sensor_data.json
$ ./load_sensor_data.sh
```

The `load_sensor_data.sh` script reads the `sensor_data.json` line by line and submits to Logstash using the `curl` command we just saw.

We have just played one day's worth of sensor readings and taken every minute from different sensors across a few geographical locations to Logstash. The Logstash data pipeline that we built earlier should have enriched and sent the data to our Elasticsearch.

It is time to switch over to Kibana and get some insights from the data.

Visualizing the data in Kibana

We have successfully set up the Logstash data pipeline and loaded some data using the pipeline into Elasticsearch. It is time to explore the data and build a dashboard that will help us gain some insights into the data.

Let's start by doing a sanity check to see if the data is loaded correctly. We can do so by going to Kibana **Dev Tools** and executing the following query:

```
GET /sensor_data-*/_search?size=0&track_total_hits=true
{
  "query": {"match_all": {}}
}
```

This query will search data across all indices matching the `sensor_data-*` pattern. There should be a good number of records in the index if the data was indexed correctly.

We will cover the following topics:

- Set up an index pattern in Kibana
- Build visualizations
- Create a dashboard using the visualizations

Let's go through each step.

Setting up an index pattern in Kibana

Before we can start building visualizations, we need to set up the index pattern for all indexes that we will potentially have for the Sensor Data Analytics application. We need to do this because our index names are dynamic. We will have one index per day, but we want to be able to create visualizations and dashboards that work across multiple indices of sensor data even when there are multiple indices. To do this, click on the **Index Patterns** link under the **Manage and Administer the Elastic Stack** section, as follows:

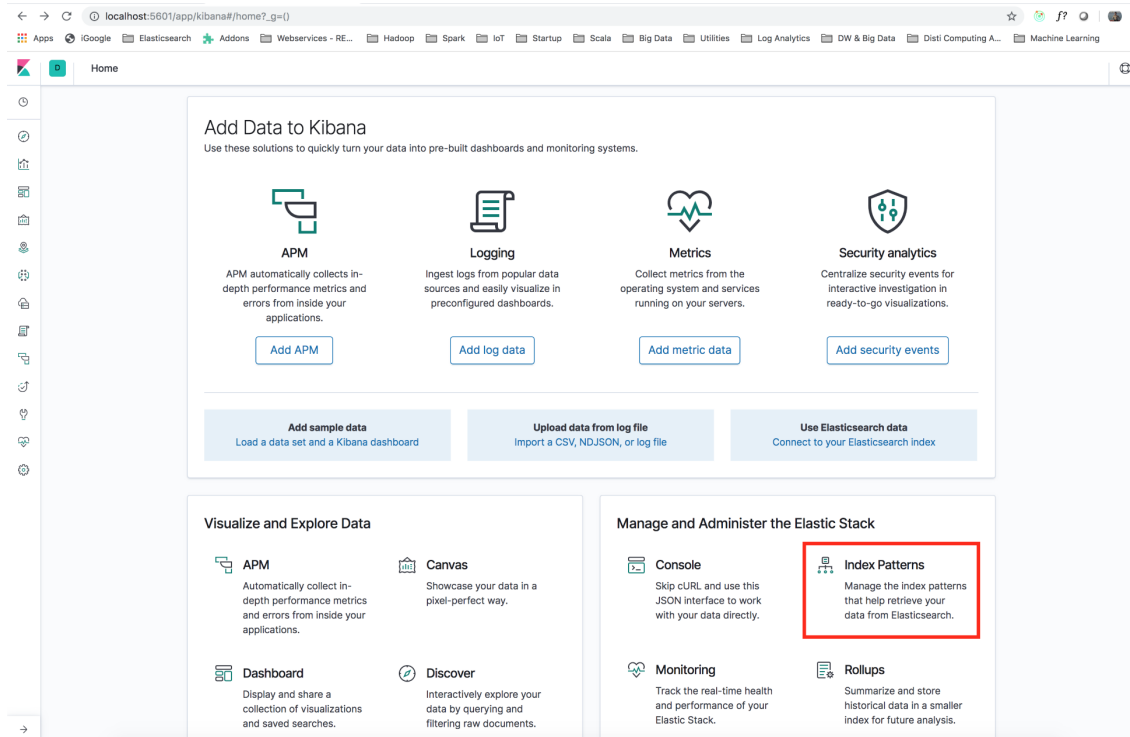


Figure 10.2: Creating an index pattern

In the **Index pattern** field, type in `sensor_data*` index pattern, as shown in the following screenshot, and click **Next step** :

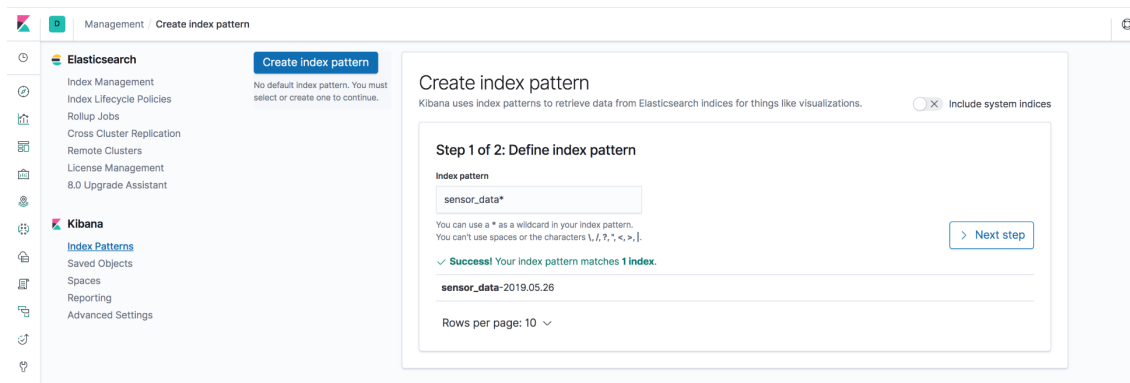


Figure 10.3: Creating an index pattern

On the next screen, in `Time Filter Field Name`, choose the **time** field as follows and click on `Create index pattern`:

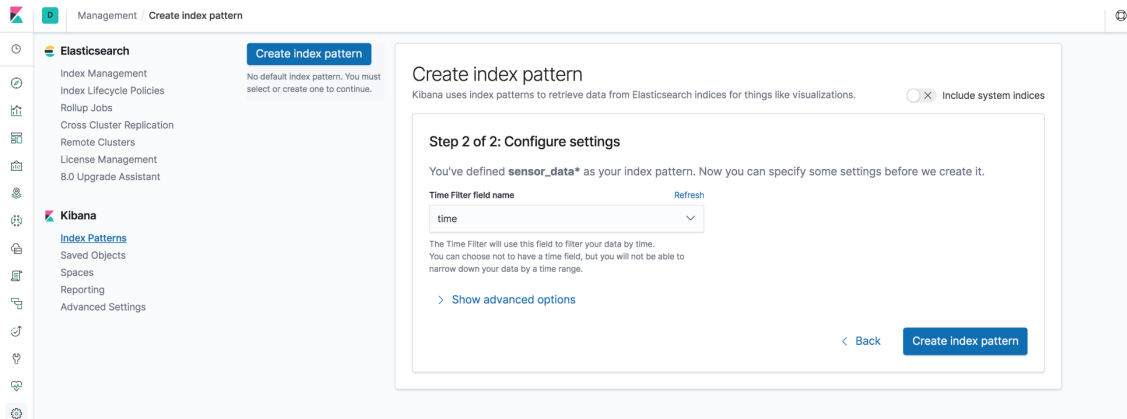


Figure 10.4: Choose Time Filter field name field for the index pattern

We have successfully created the index pattern for our sensor data. Next, we will start building some visualizations.

Building visualizations

Before we embark on an analytics project, we often already have some questions that we want to get answered quickly from visualizations. These visualizations, which answer different questions, may be packaged as a dashboard or may be used as, and when, needed. We will also start with some questions and try to build visualizations to get answers to those questions.

We will try to answer the following questions:

- How does the average temperature change over time?
- How does the average humidity change over time?
- How do temperature and humidity change at each location over time?
- Can I visualize temperature and humidity over a map?
- How are the sensors distributed across departments?

Let's build visualizations to get the answers, starting with the first question.

How does the average temperature change over time?

Here, we are just looking for an aggregate statistic. We want to know the average temperature across all temperature sensors regardless of their location or any other criteria. As we saw in Lab 7, *[Visualizing Data with Kibana]*, we should go to the `Visualize` tab to create new visualizations and click on the button with a `+ Create a Visualization` button.

Choose `Line Chart`, and then choose the `sensor_data*` index pattern as the source for the new visualization. On the next screen, to configure the line chart, follow steps 1 to 5, as shown in the following screenshot:



Figure 10.5: Creating the visualization for average temperature over time

1. Click on the time range selection fields near the top-right corner, choose **Absolute**, and select the date range as **December 1, 2017** to **December 2, 2017**. We have to do this because our simulated sensor data is from **December 1, 2017**.
2. Click on **Add a filter** as shown in Figure-10.5 and choose the **Filter** as follows:
sensorType:Temperature. Click on the **Save** button. We have two types of sensors, **Temperature** and **Humidity**. In the current visualization that we are building, we are only interested in the temperature readings. This is why we've added this filter.
3. From the **Metrics** section, choose the values shown in Figure 10.5. We are interested in the average value of the readings. We have also modified the label to be **Average Temperature**.
4. From the **Buckets** section, choose the **Date Histogram** aggregation and the **time** field, with the other options left as they are.
5. Click on the triangular **Apply changes** button.

The result is the average temperature across all temperature sensors over the selected time period. This is what we were looking for when we started building this visualization. From the preceding graph, we can quickly see that on December 1, 2017 at 15:00 IST, the temperature became unusually high. The time may be different on your machine. We may want to find out which underlying sensors reported the higher-than-normal temperatures that caused this peak.

We can click on the **Save** link at the top bar and give this visualization a name. Let's call it **Average temperature over time**. Later, we will use this visualization in a dashboard.

Let's proceed to the next question.

How does the average humidity change over time?

This question is very similar to the previous question. We can reuse the previous visualization, make a slight modification, and create another copy to answer this question. We will start by opening the first visualization, which we saved with the name **Average temperature over time**.

Execute the steps as follows to update the visualization:

1. Click on the filter with the **sensorType: Temperature** label and click on the **Edit Filter** action.
2. Change the Filter value from **Temperature** to **Humidity** and click on **Save**.
3. Modify Custom Label from **Average Temperature** to **Average Humidity** and click on the **Apply changes** button, as shown in the following screenshot.

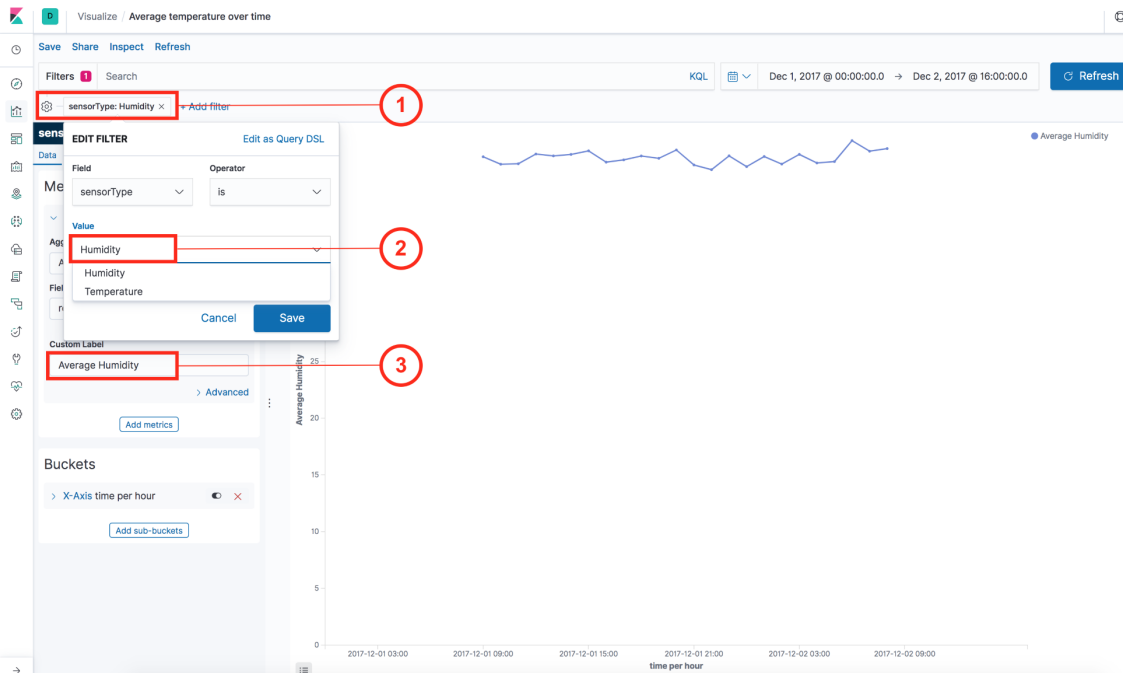


Figure 10.6: Creating the visualization for average humidity over time

As you will see, the chart gets updated for the Humidity sensors. You can click on the **Save** link at the top navigation bar. You can give a new name to the visualization, such as **Average humidity over time**, check the **Save as a new visualization** box, and click on **Save**. This completes our second visualization and answers our second question.

How do temperature and humidity change at each location over time?

This time, we are looking to get more details than the first two questions. We want to know how the temperature and humidity vary at each location over time. We will solve it for temperature.

Go to the **Visualizations** tab in Kibana and create a new **Line** chart visualization, the same as before:

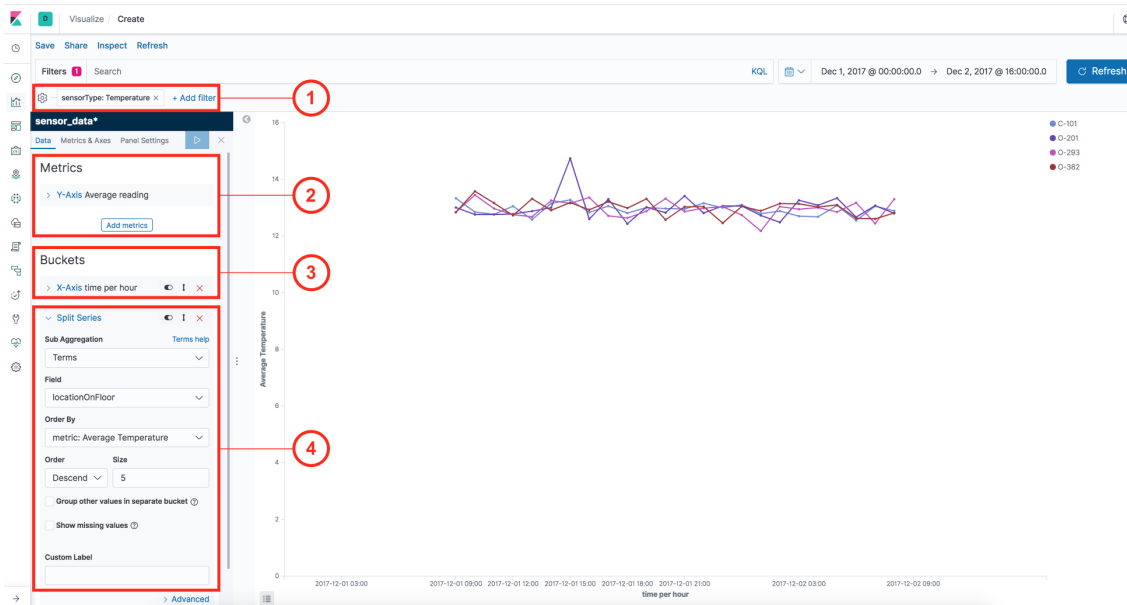


Figure 10.7: Creating the visualization for temperature at locations over time

1. Add a filter for **sensorType: Temperature** as we did before.
2. Set up the **Metrics** section exactly same as the first chart that we created, that is `Average Temperature over time` on the `reading` field.
3. Since we are aggregating the data over the **time** field, we need to choose the **Date Histogram** aggregation in the **Buckets** section. Here, we should choose the **time** field and leave the aggregation **Interval** as ****Auto**.
4. Up to this point, this visualization is the same as `Average temperature over time`. We don't just want to see the average temperature over time; we want to see it per **locationOnFloor**, which is our most fine-grained unit of identifying a location. This is why we are splitting the series using the **Terms** aggregation on the **locationOnFloor** in this step. We select **Order By** as **metric: Average Temperature**, keep **Order** as **Descend**, and **Size** to be **5** to retain only the top five locations.

We have now built a visualization that shows how the temperature changes for each value of **locationOnFloor** field in our data. You can clearly see that there is a spike in **O-201** on **December 1, 2017** at **15:00 IST**. Because of this spike, we had seen the average temperature in our first visualization spike at that time. This is an important insight that we have uncovered. Save this visualization as `Temperature at locations over time`.

A visualization for humidity can be created by following the same steps but just replacing `Temperature` with `Humidity`.

Can I visualize temperature and humidity over a map?

We can visualize temperature and humidity over the map using the **Coordinate Map** visualization. Create a new `Coordinate Map` visualization by going to the **Visualize** tab and clicking the **+** icon to create a new visualization, and perform the following steps as shown in the following screenshot:

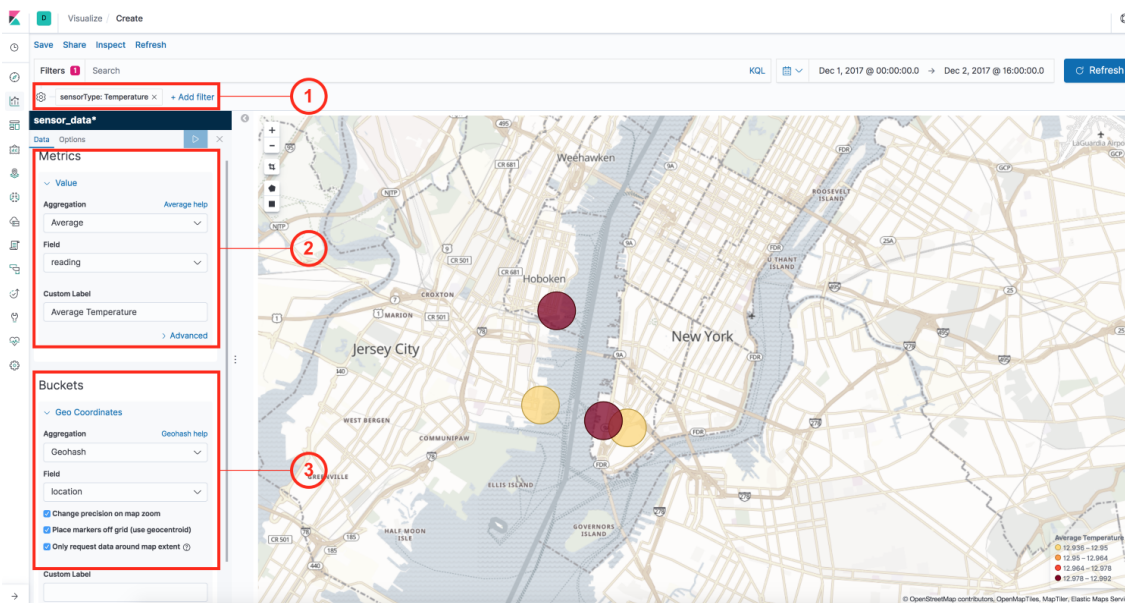


Figure 10.8: Creating a visualization to view sensor locations over a map

1. As in previous visualizations, add a filter for the **sensorType: Temperature**.
2. In the **Metrics** section, choose **Average** aggregation on the **reading** field as done previously.
3. Since this is a **Coordinate Map**, we need to choose the **GeoHash** grid aggregation and then select the **geo_point** field that we have in our data. The **location** is the field to aggregate.

As you can see, it helps in visualizing our data on the map. We can immediately see the average temperature at each site when we hover over a specific location. Focus on the relevant part of the map and save the visualization with the name `Temperature over locations`.

You can create a similar **Coordinate Map** visualization for the Humidity sensors.

How are the sensors distributed across departments?

What if we want to see how the sensors distributed across different departments? Remember, we have the `department` field in our data, which we obtained after enriching the data using the `sensor_id`. Pie charts are particularly useful to visualize how data is distributed across multiple values of a `keyword` type field, such as `department`. We will start by creating a new `pie` chart visualization.

Follow the steps as shown in the following screenshot:

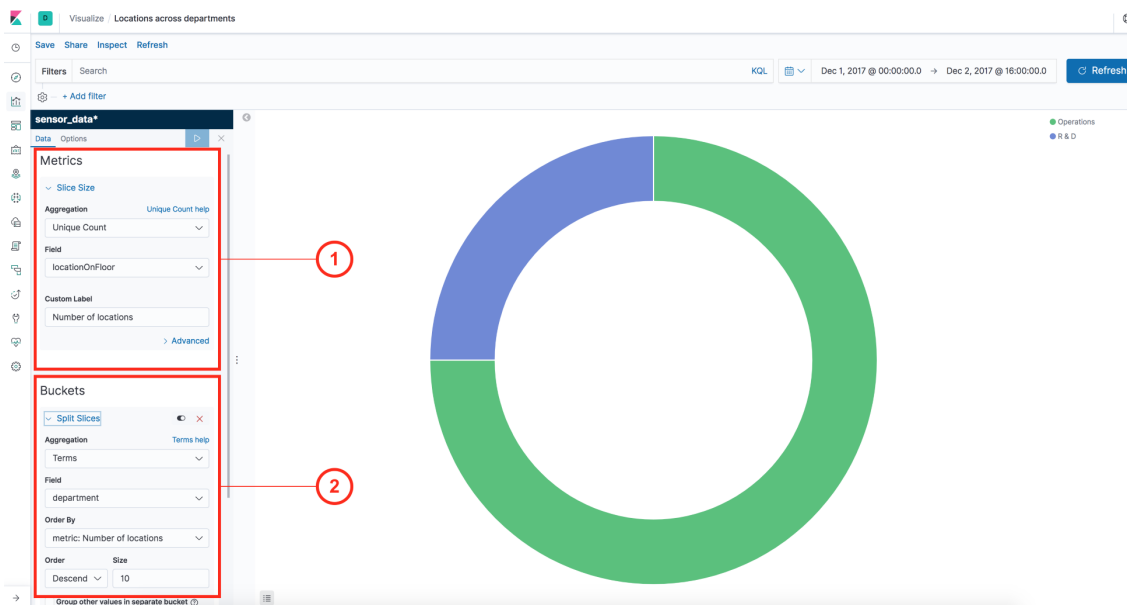


Figure 10.9: Creating a visualization for locations across departments

1. In the **Metrics** section, choose **Unique Count** aggregation and the **locationOnFloor** field. You may modify the **Custom Label** to **Number of locations**.
2. In the **Buckets** section, we need to choose **Terms** aggregation on the **department** field as we want to aggregate the data across different departments.

Click on **Apply changes** and save this visualization as **Locations across departments**. You can also create another similar visualization to visualize locations across different buildings. Let's call that visualization **Locations across buildings**. This will help us see how many locations are being monitored in each building.

Next, we will create a dashboard to bring together all the visualizations we have built.

Creating a dashboard

A dashboard lets you organize multiple visualizations together, save them, and share them with other people. The ability to look at multiple visualizations has its own benefits. You can filter the data using some criteria and all visualizations will show the data filtered by the same criteria. This ability lets you uncover some powerful insights. It can also answer more complex questions.

Let us build a dashboard from the visualizations that we have created so far. Please click on the **Dashboard** tab from the left-hand-side navigation bar in Kibana. Click on the **+ Create new dashboard** button to create a new dashboard.

Click on the **Add** link to add visualizations to your newly created dashboard. As you click, you will see all the visualizations we have built in a dropdown selection. You can add all the visualizations one by one and drag/resize to create a dashboard that suits your requirements.

Let us see what a dashboard may look like for the application that we are building:



Figure 10.10: Dashboard for sensor data analytics application

With the dashboard, you can add filters by clicking on the **Add filter** link near the top-left corner of the dashboard. The selected filter will be applied to all the charts.

The visualizations are interactive; for example, clicking on one of the pies of the donut charts will apply that filter globally. Let's see how this can be helpful.

When you click on the pie for 222 Broadway building in the donut chart at the bottom-right corner, you will see the filter for `buildingName: "222 Broadway"` added to the filters. This lets you see all of the data from the perspective of all the sensors in that building:



Figure 10.11: Interacting with the visualizations in a dashboard

Let us delete that filter by hovering over the **buildingName: "222 Broadway"** filter by clicking on the trash icon. Next, we will try to interact with one of the line charts, that is, the `Temperature at locations over`

time visualization.

As we observed earlier, there was a spike on December 1, 2017 at 15:00 IST. It is possible to zoom in to a particular time period by clicking, dragging, and drawing a rectangle around the time interval that we want to zoom in to within any line chart. In other words, just draw a rectangle around the spike, dragging your mouse while it is clicked. The result is that the time filter applied on the entire dashboard (which is displayed in the top-right corner) is changed.

Let's see whether we get any new insights from this simple operation to focus on that time period:



Figure 10.12: Zooming into a time interval from a line chart

We uncover the following facts:

1. The temperature sensor at location **O-201** (pink legend in fig-10.12) is steadily rising around this time.
2. In the **Coordinate Map** visualization, you can see that the highlighted circle is red, compared to the other locations, which are yellow. This highlights that the location has an abnormally high temperature compared to the other locations.

Interacting with charts and applying different filters can provide powerful insights like the ones we just saw.

This concludes our application and demonstration of what we can do using the Elastic Stack components.

Summary

In this lab, we built a sensor data analytics application that has a wide variety of applications, as it is related to the emerging field of IoT. We understood the problem domain and the data model, including metadata related to sensors. We wanted to build an analytics application using only the components of the Elastic Stack, without using any other tools and programming languages, to obtain a powerful tool that can handle large volumes of data.